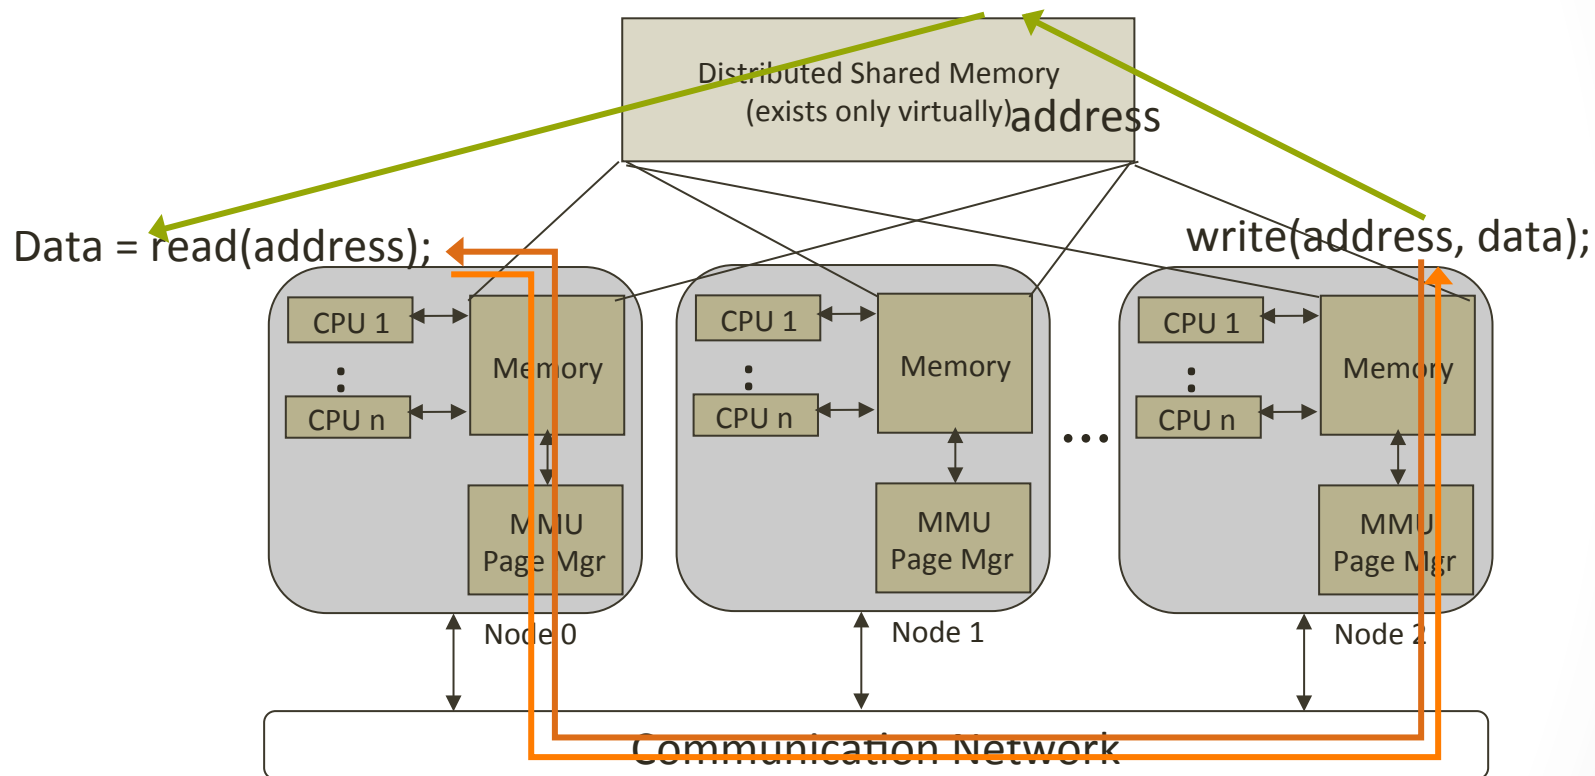


Distributed Shared Memory (DSM)

17-Nov-2015

DSM Basics



A cache line or a page is transferred to and cached in the requested computer.

Simple example

```
struct shared { int a, b; };
```

Program Writer:

```
main()
{
    struct shared *p = (struct shared *) allocShared(...)
    p->a = p->b = 0;           /* initialize fields to zero */
    while(TRUE) {             /* continuously update structure fields */
        p->a = p->a + 1;
        p->b = p->b - 1;
    }
}
```

Program Reader:

```
main()
{
    struct shared *p = ...
    while(TRUE) { /* read the fields once every second */
        printf("a = %d, b = %d\n", p->a, p->b);
        sleep(1);
    }
}
```

Why DSM

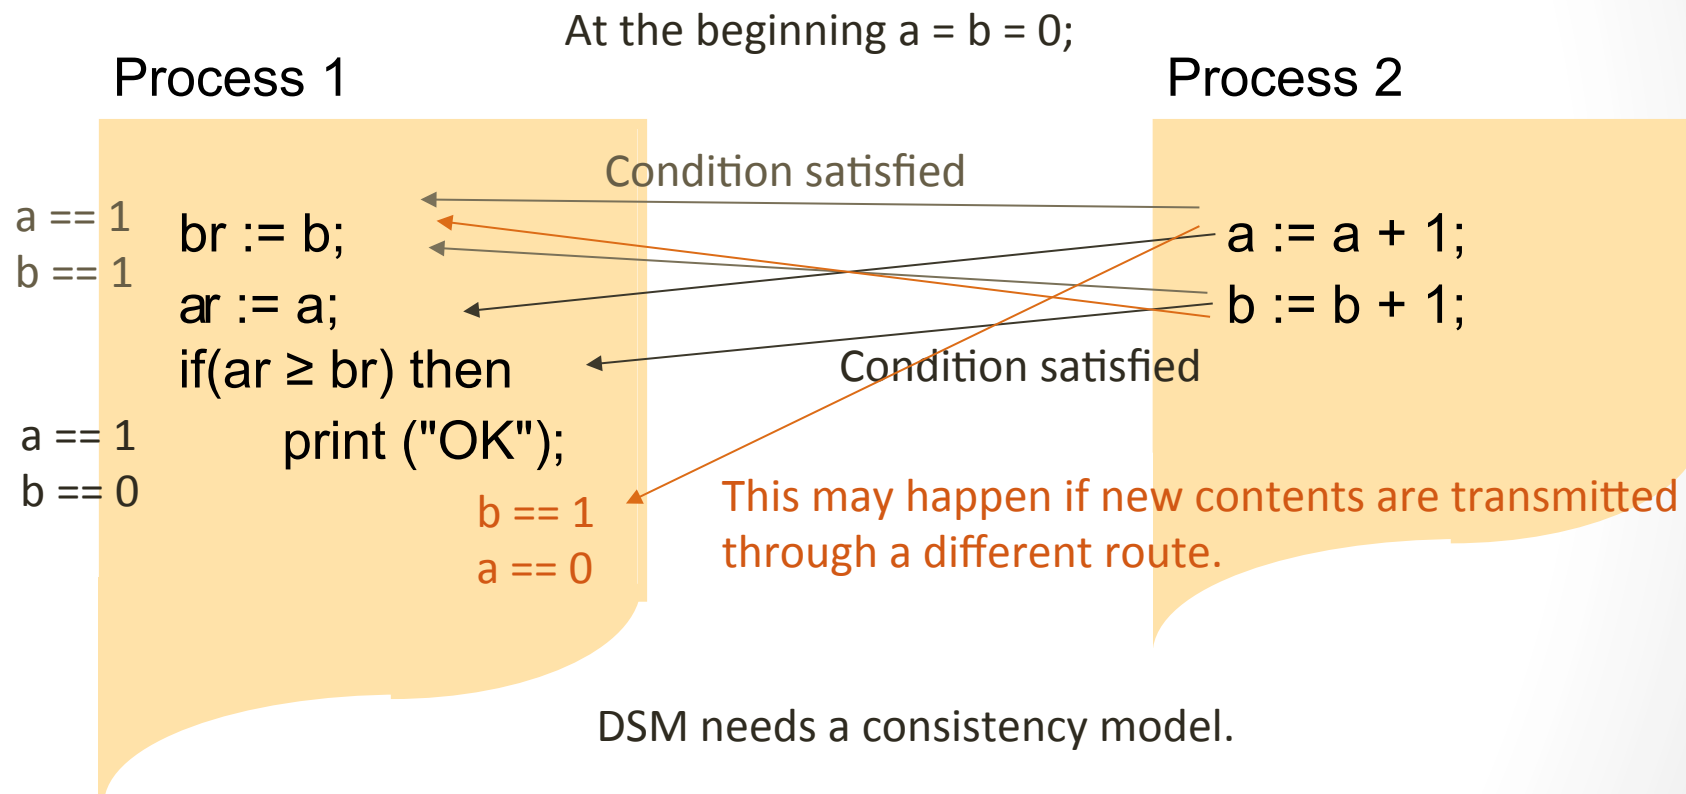
- Simpler abstraction
 - Underlying tedious communication primitives shielded by memory accesses
- Better portability of distributed programs
 - Natural transition from sequential to distributed application
- Better performance of some applications
 - Data locality, one-demand data movement, and large memory space reduce network traffic and paging/swapping activities.
- Flexible communication environment
 - Sender and receiver have no need to know each other. They even need not coexist.
- Ease of process migration
 - Migration is completed only by transferring the corresponding PCB to the destination.

Main Issues

- Granularity
 - **Fine** (less false sharing but more network traffic)
 - Cache line (e.g. Dash and Alewife), Object (e.g. Orca and Linda), Page (e.g. Ivy)
 - **Coarse**(more false sharing but less network traffic)
- Memory consistence and access synchronization
 - Strict, Sequential, Causal, Weak, and Release Consistency models
- Data location and access
 - Broadcasting, centralized data locator, fixed distributed data locator, and dynamic distributed data locator
- Replacement strategy
 - LRU or FIFO (The same issue as OS virtual memory)
- Thrashing
 - How to prevent a block from being exchanged back and forth between two nodes.
- Heterogeneity
- Implementation
 - hardware implementation, OS implementation, and User-level implementation.

Consistency Models

Two processes accessing shared variables




Consistency Models


Strict Consistency

- $W_i(x, a)$: Processor i writes a on variable x , (i.e., $x = a$);).
- $b \leftarrow R_i(x)$: Processor i reads b from variable x . (i.e., $y = x$; && $y == b$);).
- Any read on x must return the value of the most recent write on x .

Strict Consistency




Not Strict Consistency



Consistency Models

Linearizability and Sequential Consistency

- **Linearizability:** Operations of each individual process appear to all processes in the same order as they happen.
- **Sequential Consistency:** Operations of each individual process appear in the same order to all processes.




Consistency Models


FIFO and Processor Consistency

- **FIFO Consistency:** writes by a single process are visible to all other processes in the order in which they were issued.
- **Processor Consistency:** FIFO Consistency + all write to the same memory location must be visible in the same order.

FIFO Consistency



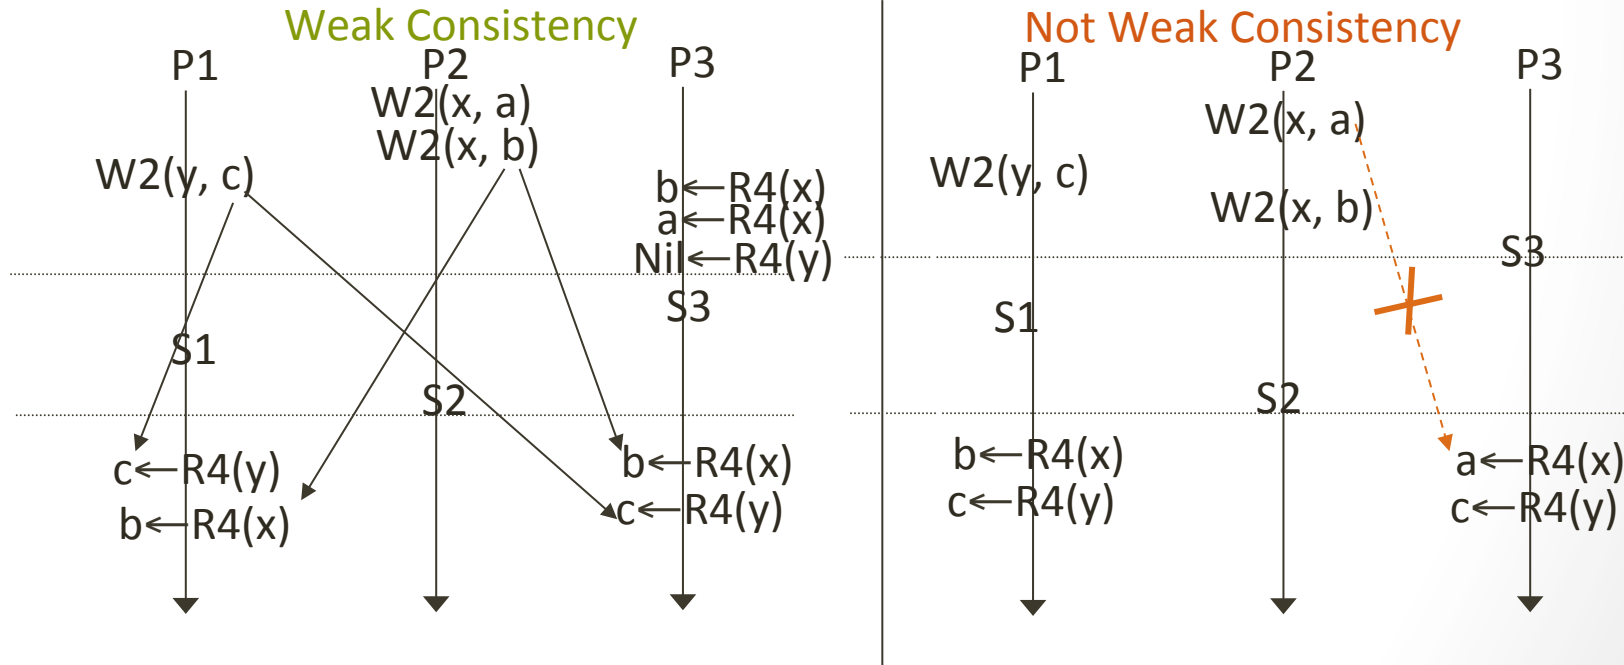
Processor Consistency



Consistency Models

Weak Consistency


- Accesses to synchronization variables must obey sequential consistency.
- All previous writes must be completed before an access to a synchronization variable.
- All previous accesses to synchronization variables must be completed before access to non-synchronization variable.



Consistency Models

Release Consistency

- Access to acquire and release variables obey processor consistency.
- Previous acquires requested by a process must be completed before the process performs a data access.
- All previous data accesses performed by a process must be completed before the process performs a release.



Consistency Models

Release Consistency (Example)

Process 1:


```
acquireLock();           // enter critical section  
a := a + 1;  
b := b + 1;  
releaseLock();          // leave critical section
```

Process 2:

```
acquireLock();           // enter critical section  
print ("The values of a and b are: ", a, b);  
releaseLock();          // leave critical section
```

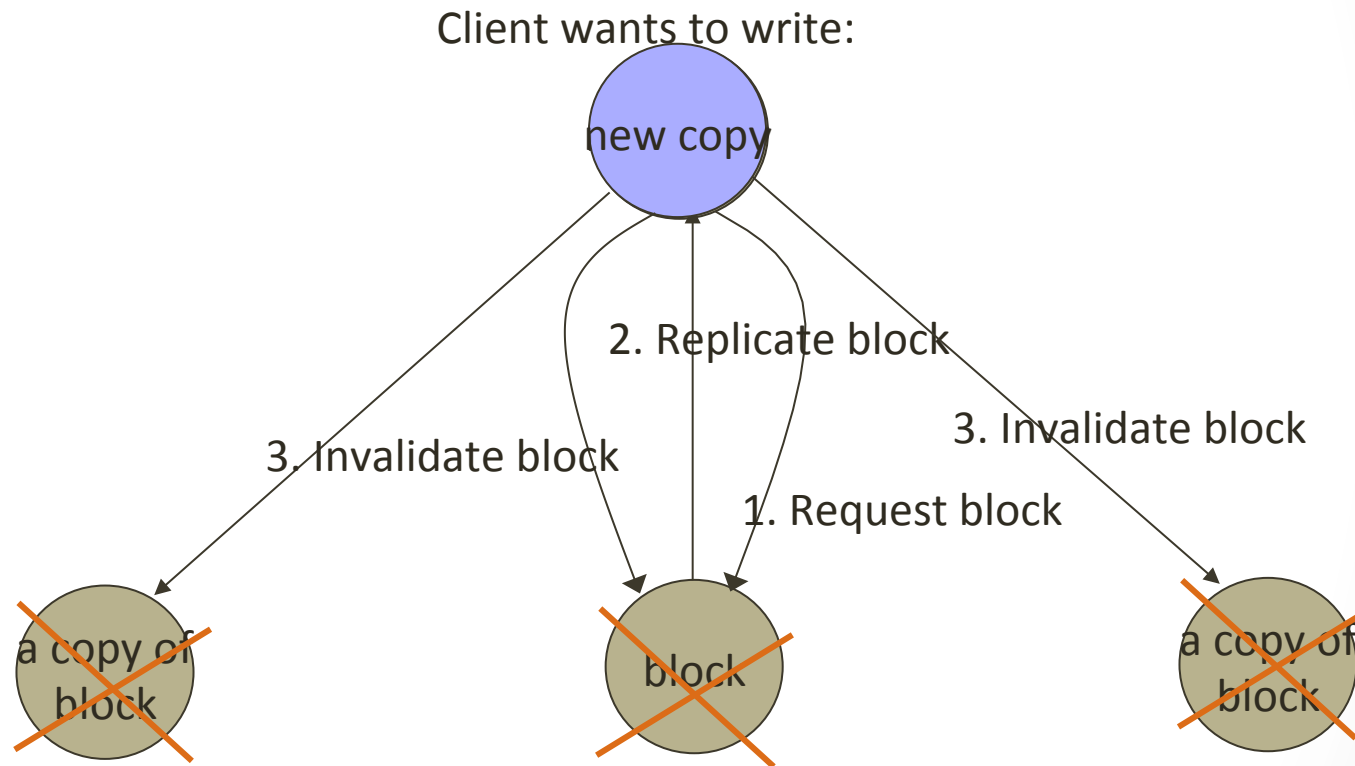
Implementing Sequential Consistency

Replicated and Migrating Data Blocks



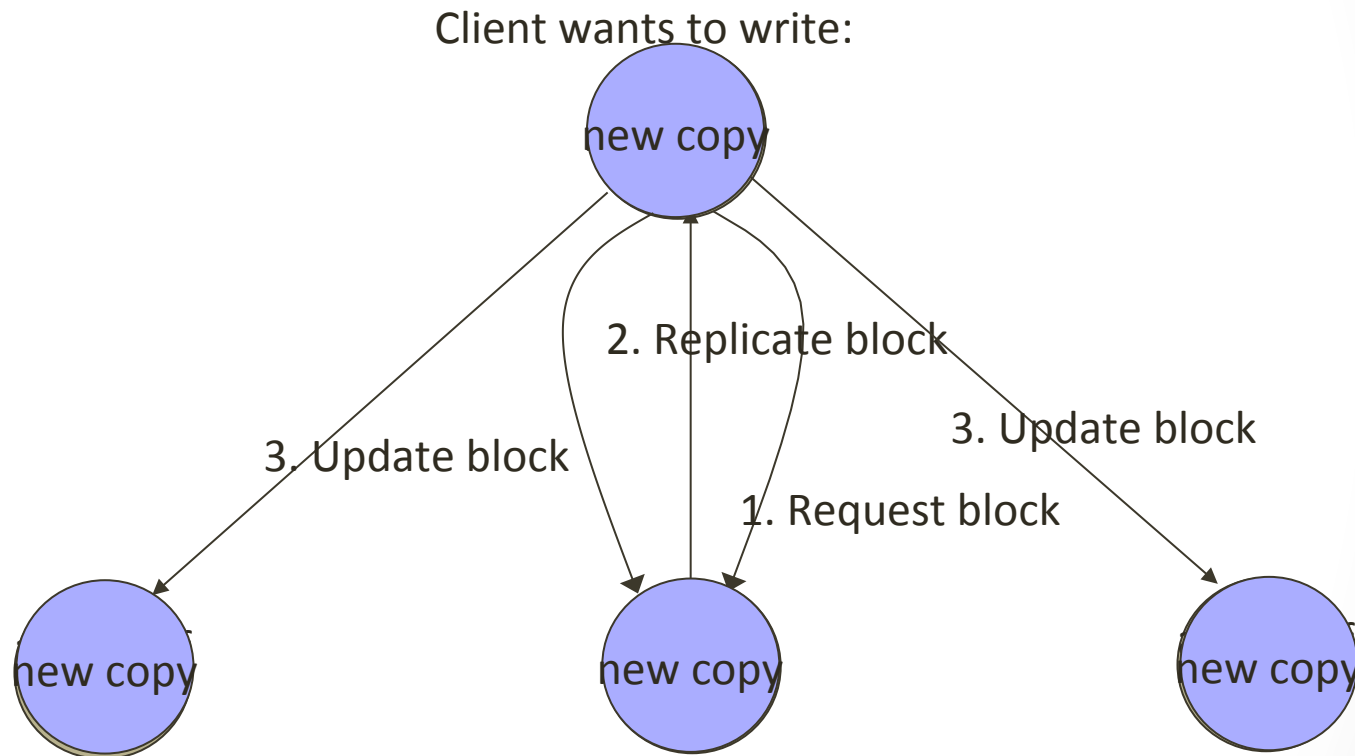
Implementing Sequential Consistency

Write Invalidation




Implementing Sequential Consistency

Write Update




Implementing Sequential Consistency

Read/Write Request



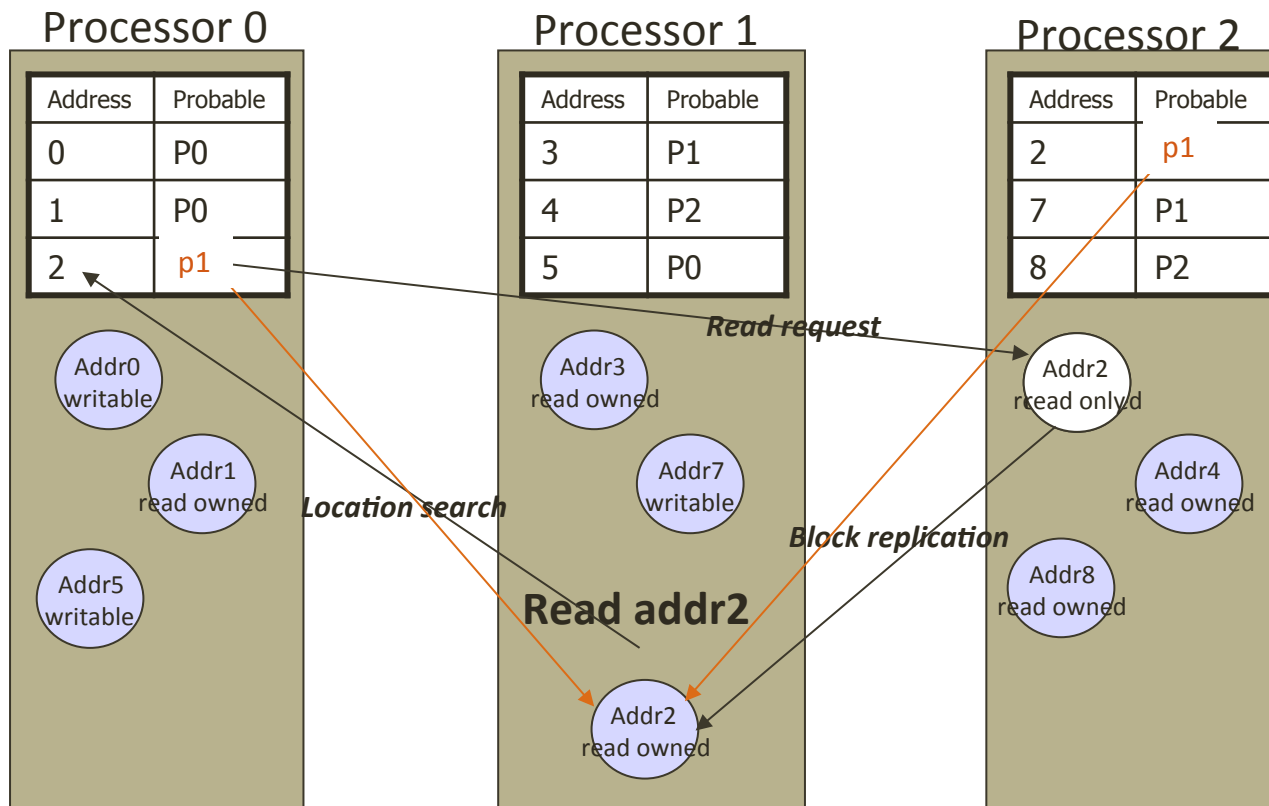
Implementing Sequential Consistency

Locating Data –Fixed Distributed-Server Algorithms



Implementing Sequential Consistency

Locating Data – Dynamic Distributed-Server Algorithms



- Breaking the chain of nodes:
 - When the node receives an invalidation
 - When the node relinquishes ownership
 - When the node forwards a fault request
- The node points to a new owner

Replacement Strategy

- Which block to replace
 - Non-usage based (e.g. FIFO)
 - Usage based (e.g. LRU)
 - Mixed of those (e.g. Ivy)
 - Unused/Nil: replaced with the highest priority
 - Read-only: the second priority
 - Read-owned: the third priority
 - Writable: the lowest priority and LRU used.
- Where to place a replaced block
 - Invalidating a block if other nodes have a copy.
 - Using secondary store
 - Using the memory space of other nodes

Thrashing

- Thrashing:
 - Two or more processes try to write the same shared block.
 - An owner keeps writing its block shared by two or more reader processes.
 - The larger a block, the more chances of false sharing that causes thrashing.
- Solutions:
 - Allow a process to prevent a block from accessed from the others, using a lock.
 - Allow a process to hold a block for a certain amount of time.
 - Apply a different coherence algorithm to each block.
- What do those solutions require users to do?
- Are there any perfect solutions?